

Resolving Conflicting Constraints in Multi-Agent Reinforcement Learning with Layered Safety

Jason J. Choi*, Jasmine Jerry Aloor*, Jingqi Li*, Maria G. Mendoza,
Hamsa Balakrishnan, Claire J. Tomlin

Collision-free operation is a fundamental requirement

An aerial photograph of a roundabout with a central island covered in pink and red flowers. Several cars are visible on the road, including a dark blue car at the top, a red car at the top right, a white car at the bottom right, and a black car at the bottom left. Pedestrians are walking on the sidewalks, and a few cyclists are on a path to the right. The road has white dashed lines and arrows indicating the flow of traffic.

for multi-robot coordination tasks
formation control,
multi-robot payload transport,
and autonomous navigation.

Growing interest in AAM applications: scalable low-altitude air traffic management system



Compared to current aviation, AAM operations are expected to be large-scale, ad hoc, on-demand, and dynamic. These characteristics motivate the development of a new air traffic management (ATM) framework that can achieve scalable, efficient, and collision-free operations

A lot of safety measures are being practiced in the current aviation industry to maintain safety.




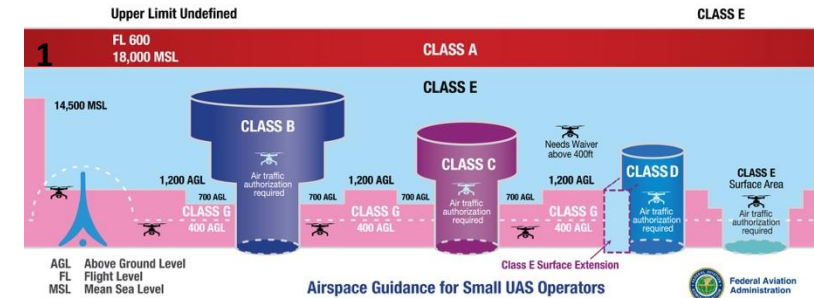
Previous work: Deconfliction & Collision Avoidance for Aviation

Airspace Separation¹

Centralized manual Air Traffic Control (ATC)²

Protocols & Conventional Practice³ built around assumptions that
at most two vehicles engage in deconfliction

Imminent collision avoidance autopilot (TCAS, ACAS⁴)
(20~30 sec before collision)



1 FAA, Airspace 101 – Rules of the Sky

3 Right-of-way rules, FAR 91.113

4 Kochenderfer et al., Lincoln Lab, 2012


Era of Advanced Aerial Mobility: New requirements

Airspace Separation

Decentralized Air Traffic Operation¹

Protocols & Conventional Practice built around assumptions that
~~at most two vehicles engage in deconfliction~~

Imminent collision avoidance autopilot (TCAS, ACAS)
(20~30 sec before collision)



⁶
1 NASA Ames, Provider of Services to UAM, 2021

Contributions: Layered

1. Layered Safety Architecture: Integration of CBV reinforcement learning, creating a two-layer system

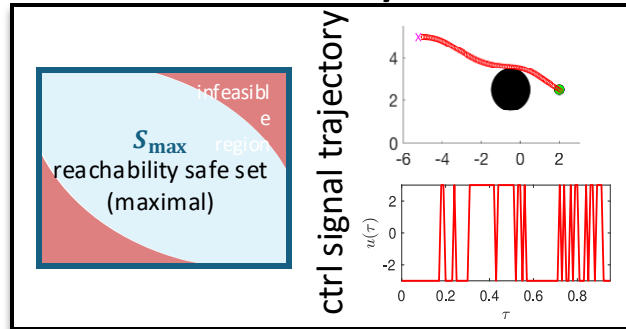
2. Safe MARL Training Method: Train using curriculum of conflicting constraint zones, without relying on per

3. Real-World & Simulated Validation: Demonstrate effectiveness through hardware experiments with Crazyflie drones and simulations in dense AAM scenarios.

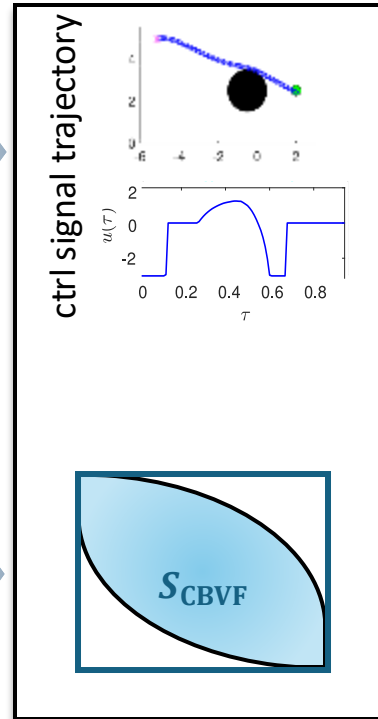


Background: **Control Barrier-Value Function (CBVF)** for Safety of Pairwise Interaction



HJ Reachability



CBVF



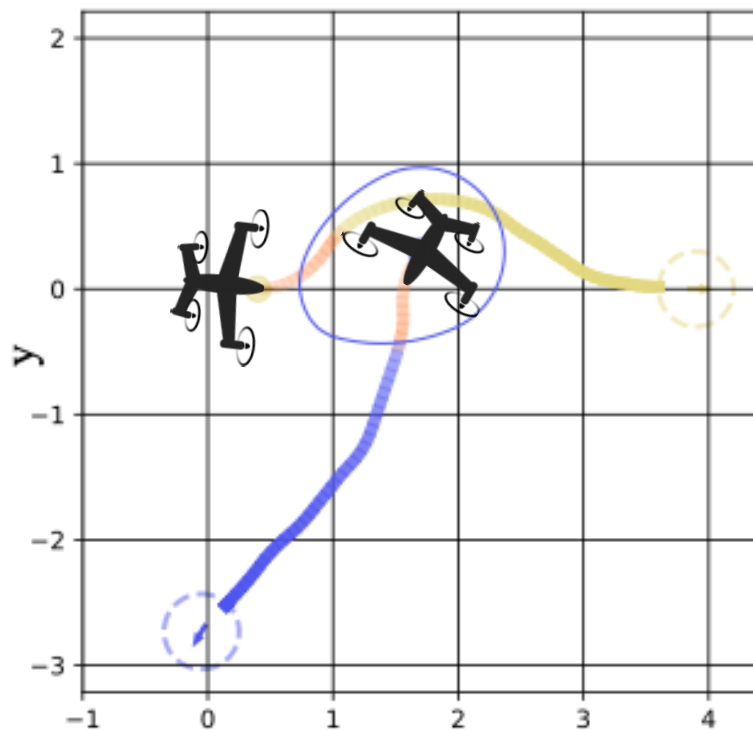
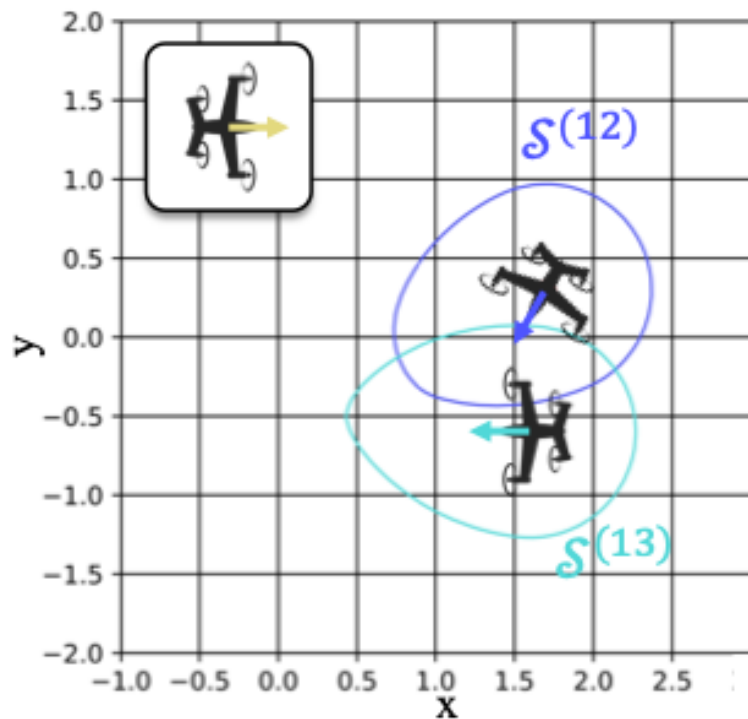
Control Barrier Functions



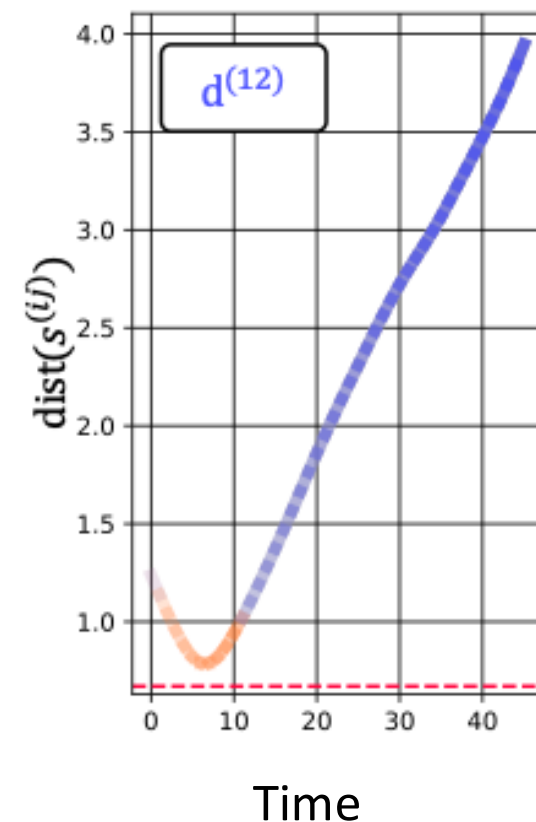
How do existing approaches work?

Traditional Control: Safety-guaranteed, but struggles with complex multi-agent interactions.

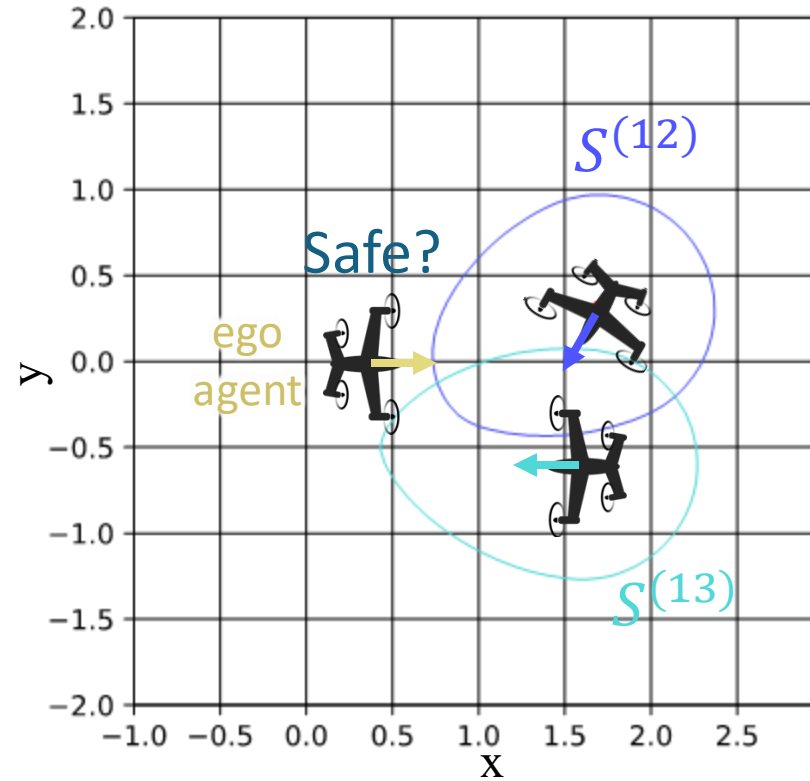
E.g., Two agents



Inter-aircraft distance




Now with three agents

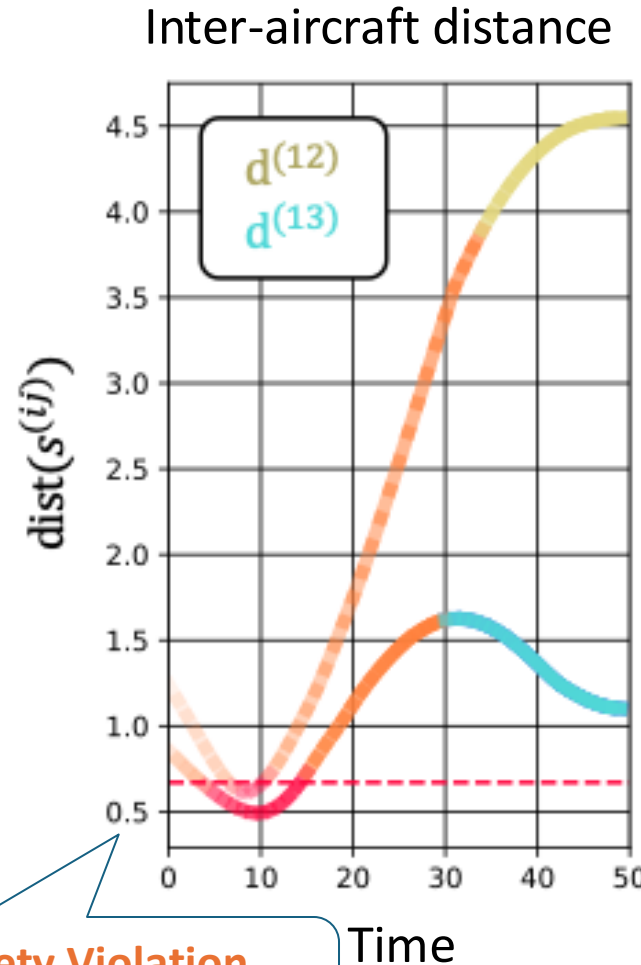


  : Max safe set boundaries from CBVF computation

Resolving a constraint with two agents may conflict with others



Safety Violation
(Near-Collision)



- Creates a gridlock
- Inability to satisfy all constraints
- Collision possible

Key Challenge: Intersection of individual safe sets is not the true combined safe sets

Aka "Leaky corner"

$$S(C_1) \cap S(C_2) \neq S(C_1 \cap C_2)$$


Distance Ranges Based On Potential Conflict Range

Definition (informal):

If more than two vehicles are within the **potential conflict zone**, the leaky corner issue can arise and the safety of all agents is not guaranteed anymore.

2. r_{conflict} : Potential conflict range

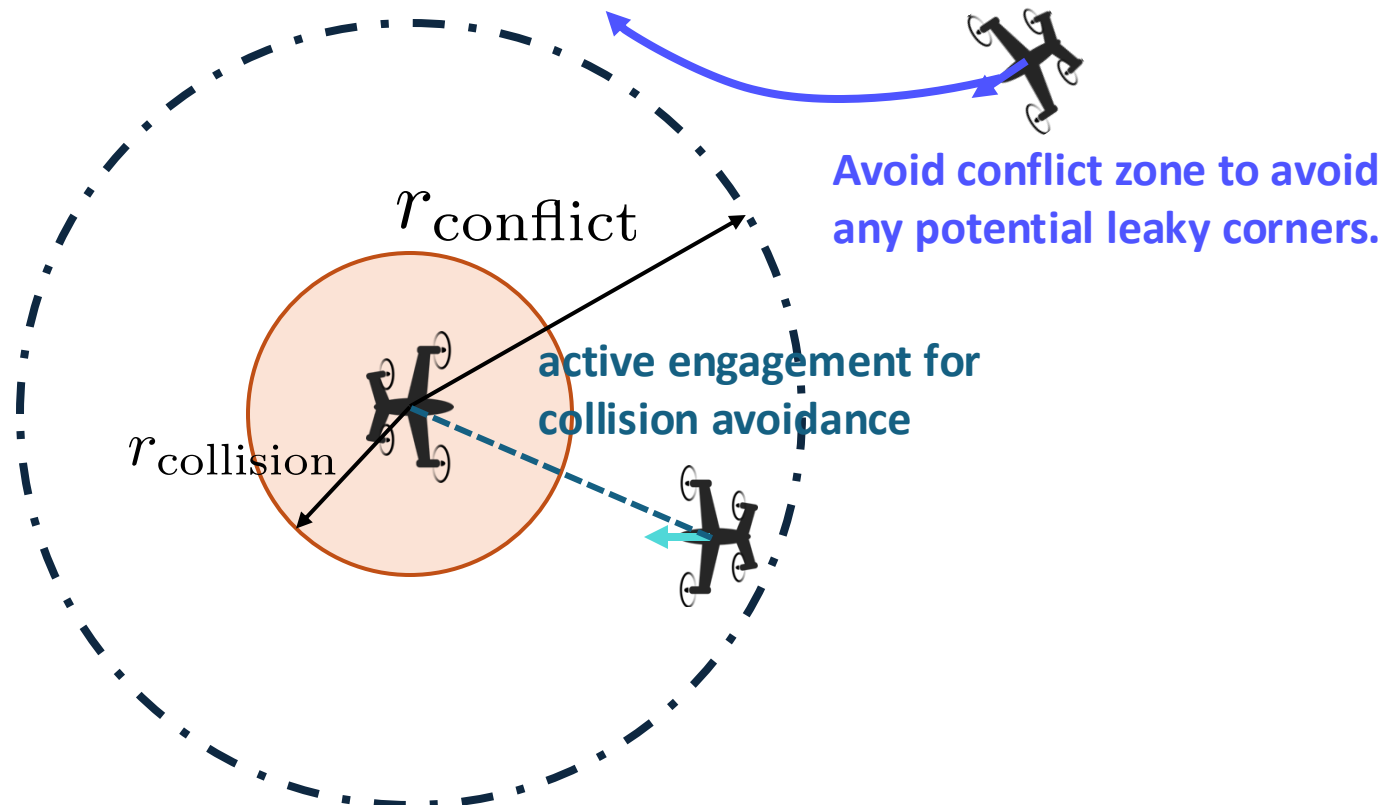
1. $r_{\text{collision}}$: Where collision is imminent



Distance Ranges Based On Potential Conflict Range

Definition (informal):

If more than two vehicles are within the **potential conflict zone**, the leaky corner issue can arise and the safety of all agents is not guaranteed anymore.




Safety-Informed Multi-Agent RL (MARL)

Layered Safety for MARL:

Safe-informed
Graphical Neural Network
MARL policy

1. Based on the agent's local observations, MARL Policy learns to **navigate to waypoints while avoiding the potential conflict regions.**



Safety-Informed Multi-Agent RL (MARL)


Layered Safety for MARL:

Safe-informed
Graphical Neural Network
MARL policy




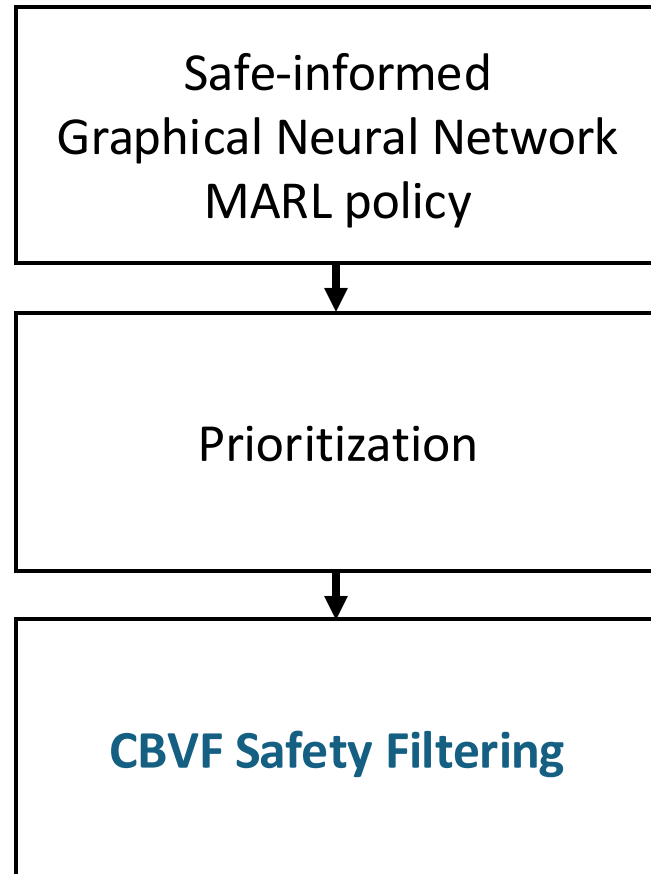
Prioritization

2. Evaluate CBVF values against each agents
and pick the agent with **worst CBVF value**.



Safety-Informed Multi-Agent RL (MARL)

Layered Safety for MARL:



Investigating Varying Levels of Safety

1. Policy trained without the safety filter and no safety penalty (**safety blind**)
2. Policy trained with the safety filter and no safety penalty (no penalty).
3. Policy trained with the safety filter and with r_{conflict} violation penalty (**proposed**).

Low safety

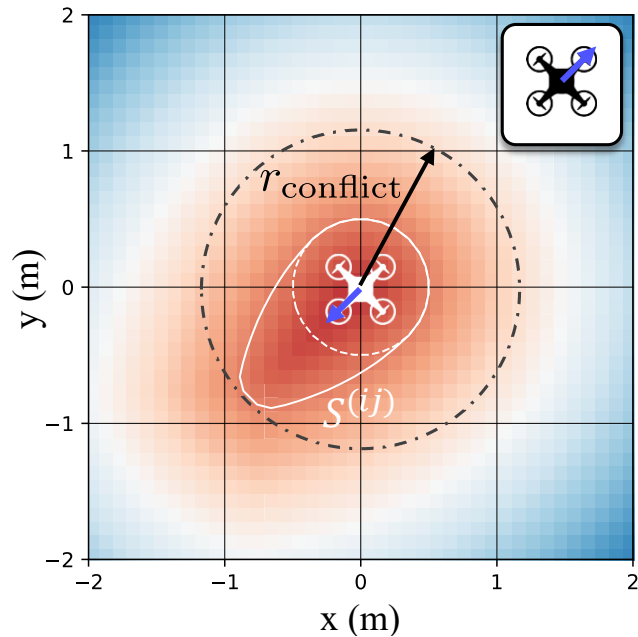


High safety

Conflict zone & safe sets computed with CBVF


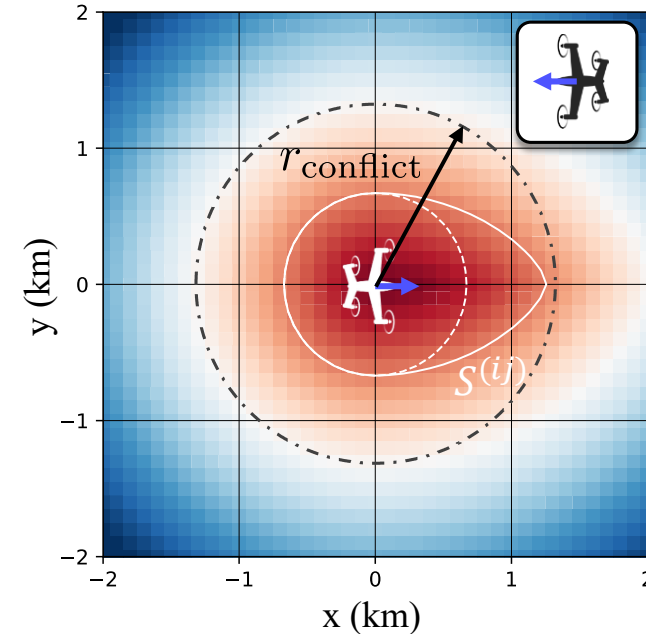
Quadrotor:

can hover & hold position for safety

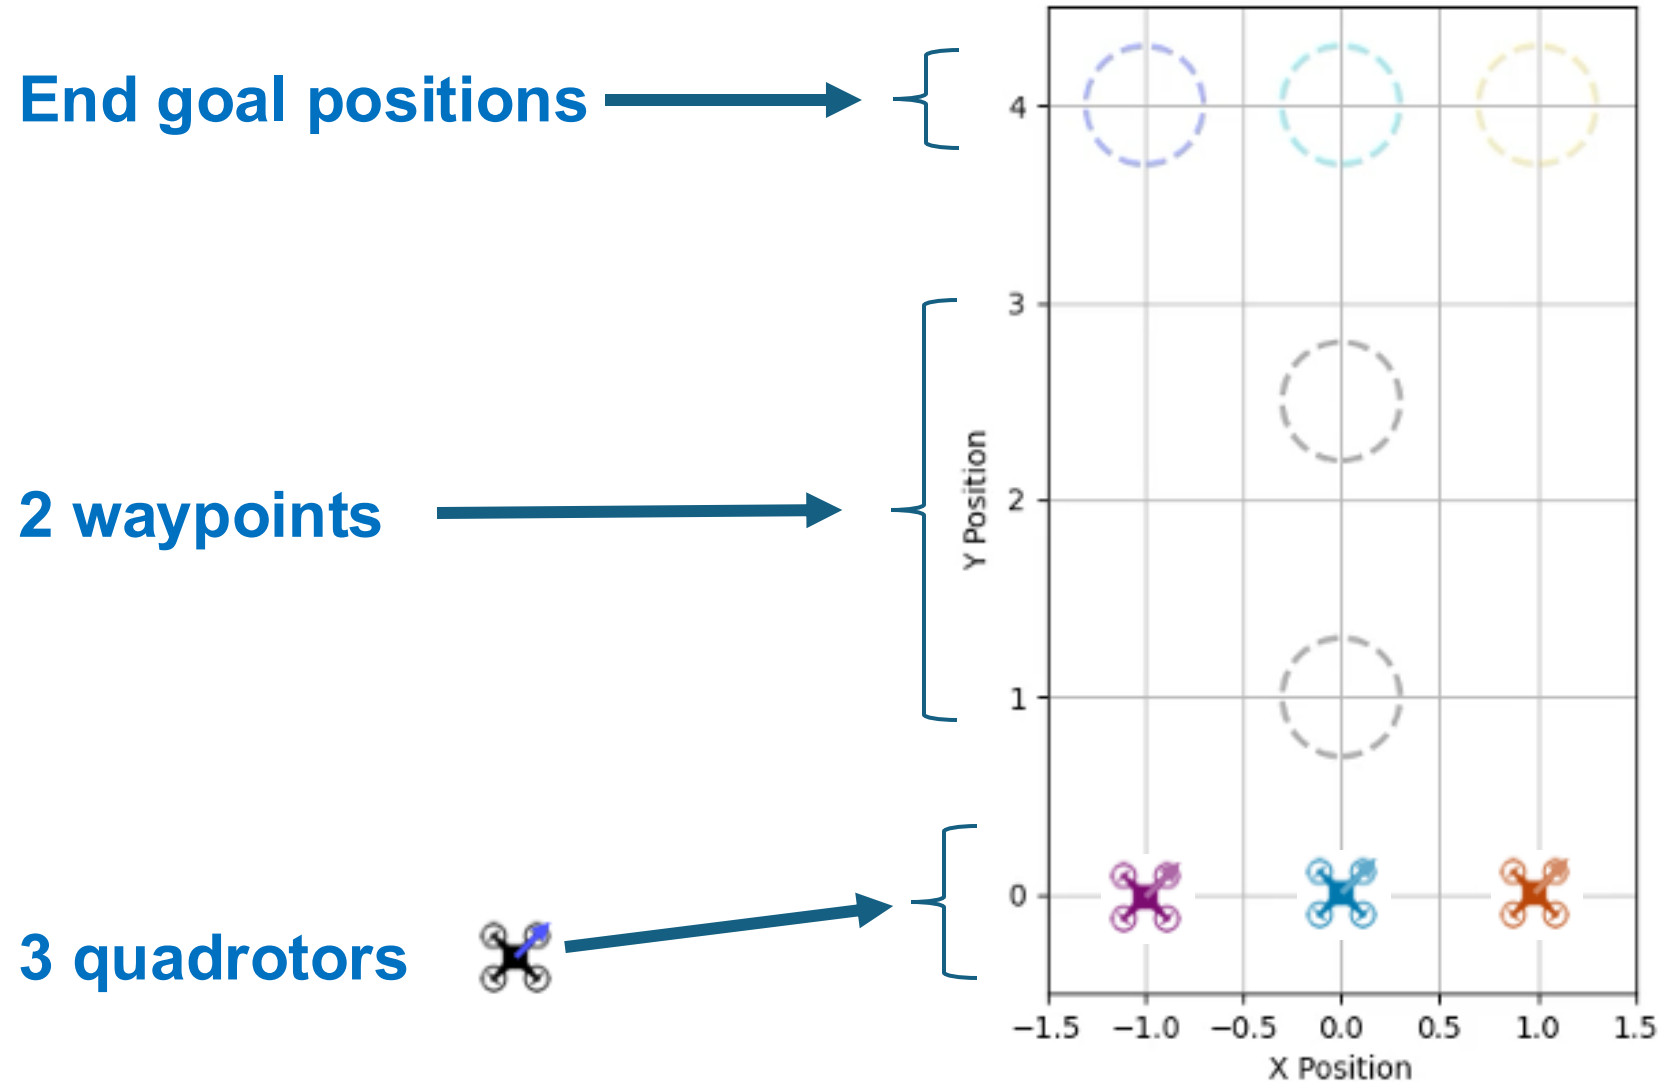


Wingborne Air Taxi:

Needs active turn for safety ($v_{\min} > 0$)

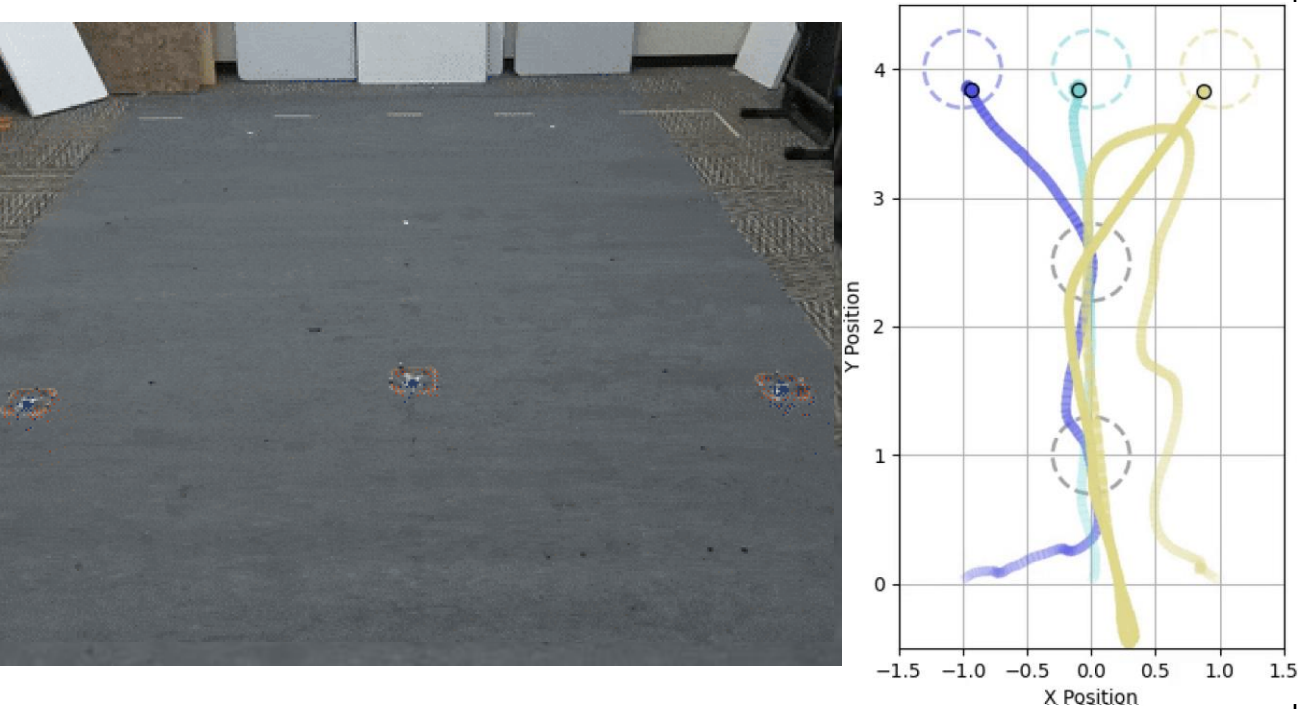


Safe Navigation Of Quadrotors Through Waypoints

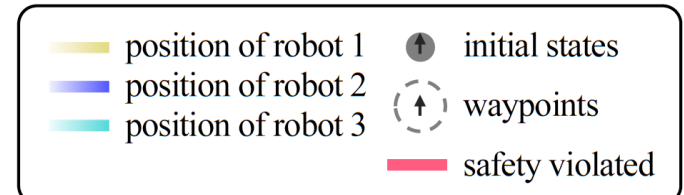
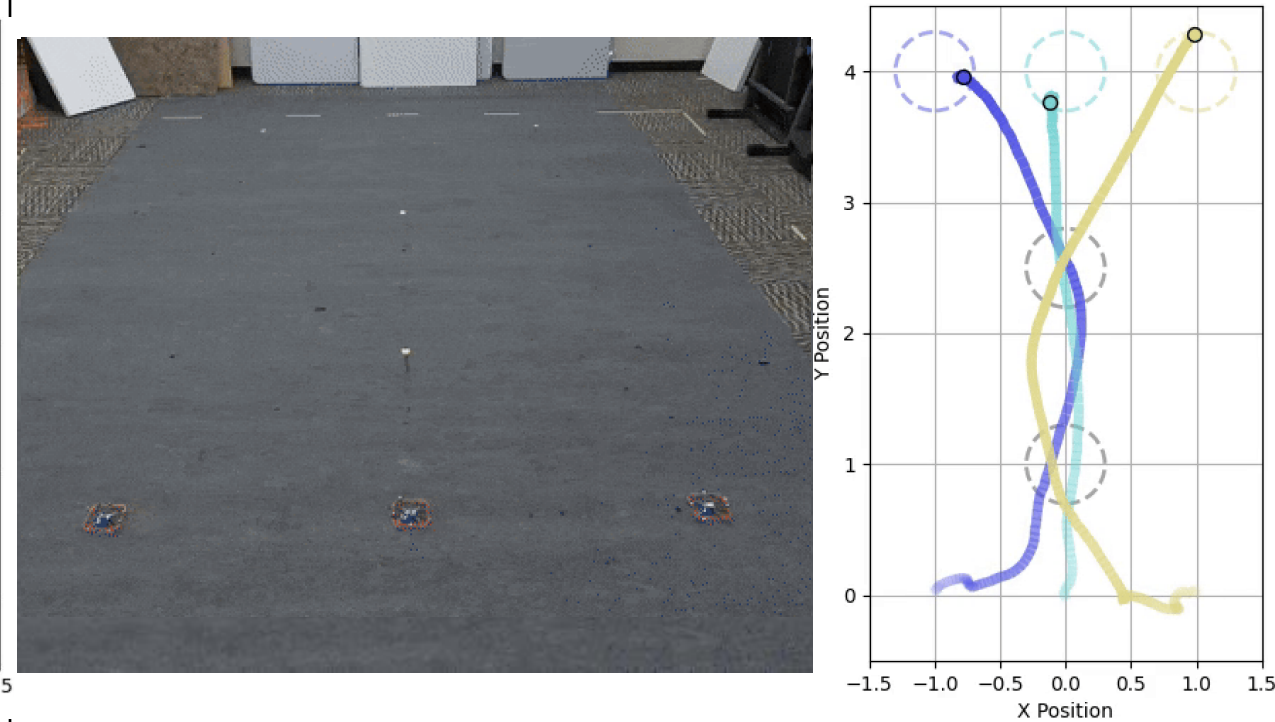


Safe Navigation Of Quadrotors Through Waypoints

Safety-blind MARL (safety filtered)



Safety-informed MARL



2x Speed

Comparisons to model-based and model-free methods

Performance: Goal reach rate (%)

Safety: Near collision (%)

Quadrotor dynamics: N=4 agents

Methods	Goal reach(%)	Near collision(%)
DG-PPO	96 ± 11.8	0.04 ± 0.16
Exponential CBF	100 ± 0	0.0 ± 0.0
Our Method	100 ± 0	0.0 ± 0.0

Quadrotor dynamics: N=8 agents

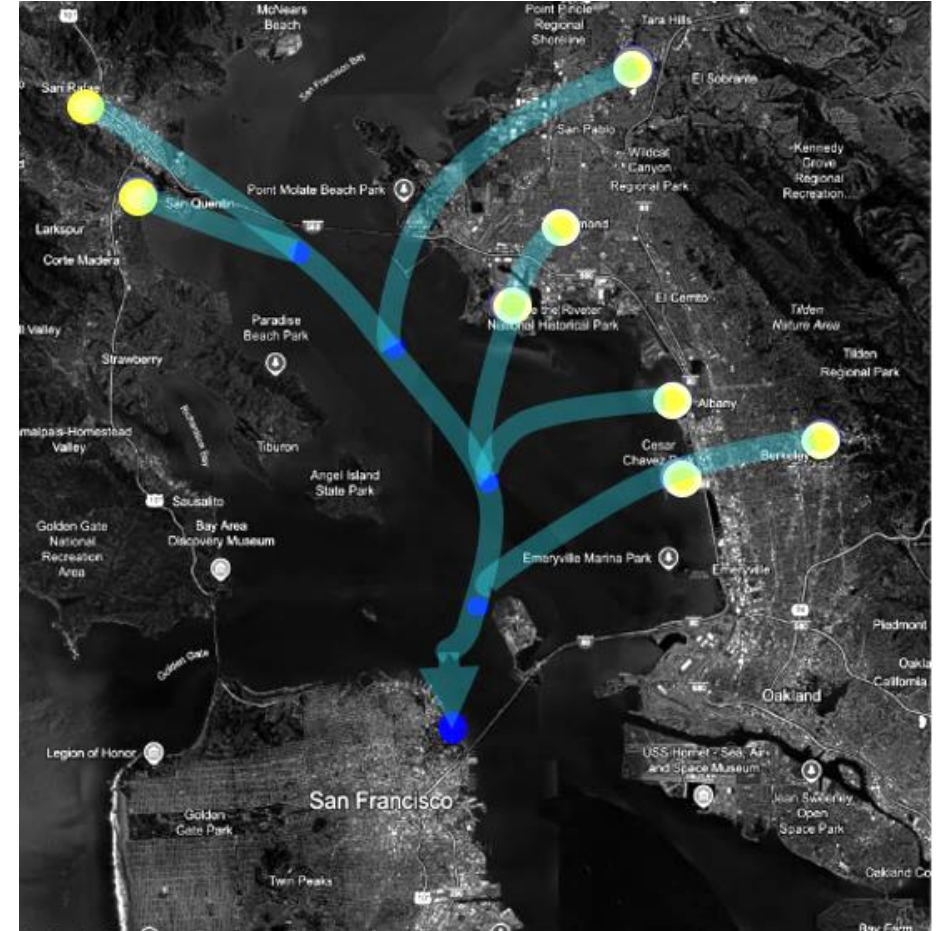
Methods	Goal reach(%)	Near collision(%)
DG-PPO	100 ± 0	9.1 ± 2.7
Exponential CBF	93 ± 8.9	8.8 ± 10.7
Our Method	100 ± 0	0.0 ± 0.0

Autonomous Urban Air Mobility Traffic Management: Air Taxi Operation Scenario

Urban Air Mobility (UAM) demand analysis

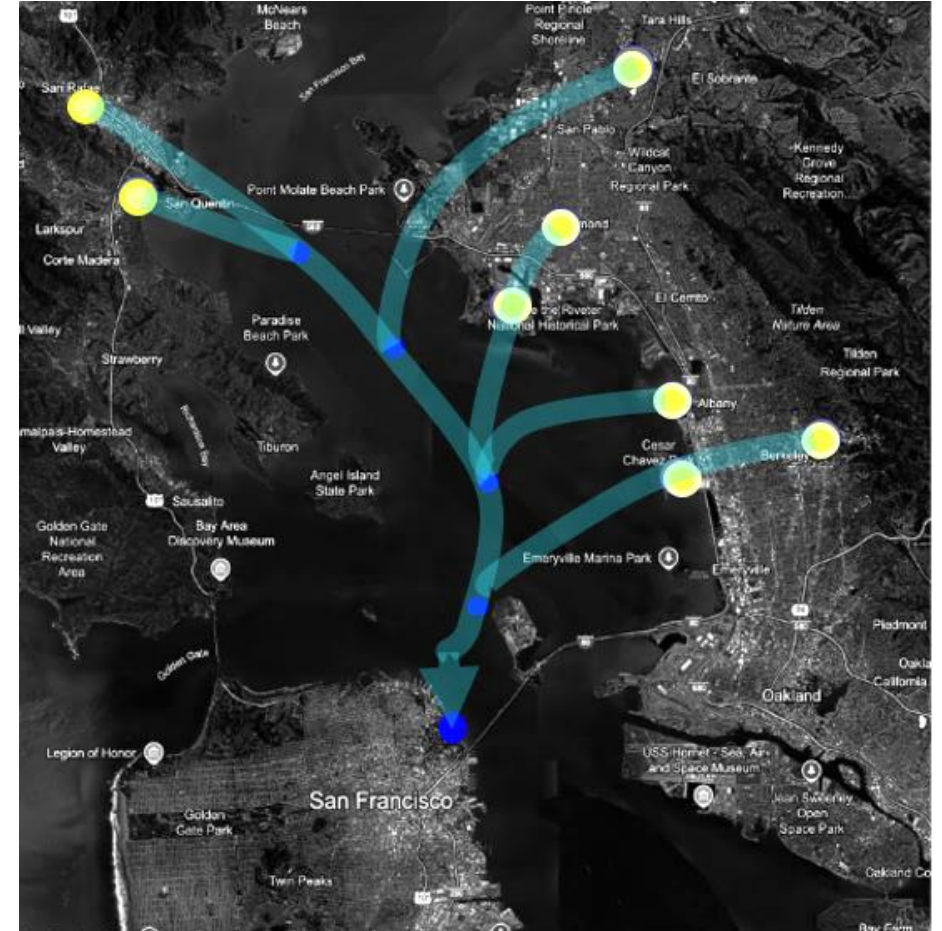
Consider a peak-density scenario:

- Each vertiport serves 125 aircraft/hour
- Two operations (takeoffs and landings)/min,
- Corridors with 1500 ft separation



(8 vehicles merging into the air corridor that
lands at San Francisco)


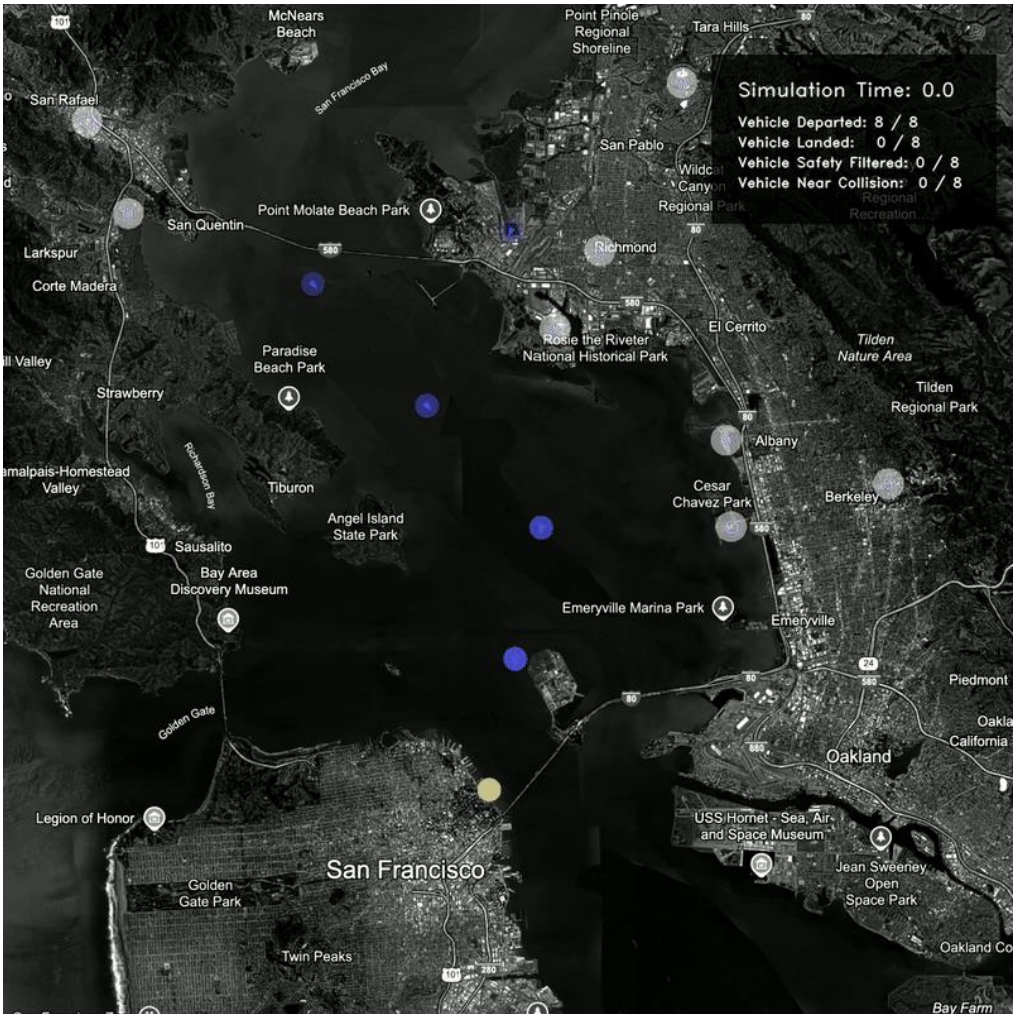
Autonomous Urban Air Mobility Traffic Management: Air Taxi Operation Scenario



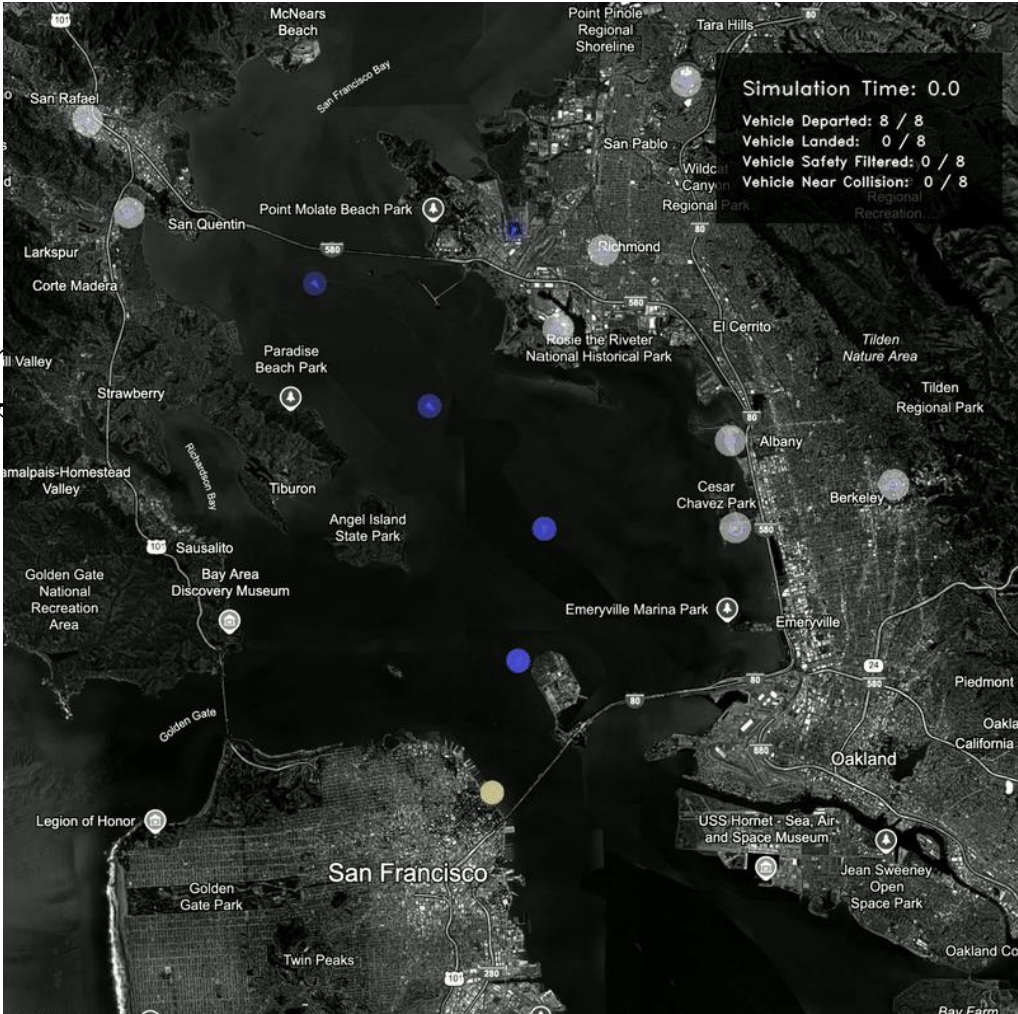
(8 vehicles merging into the air corridor that
lands at San Francisco)

Air Taxi Operation Scenario: Merging into corridor

Safety-blind MARL (safety filtered)



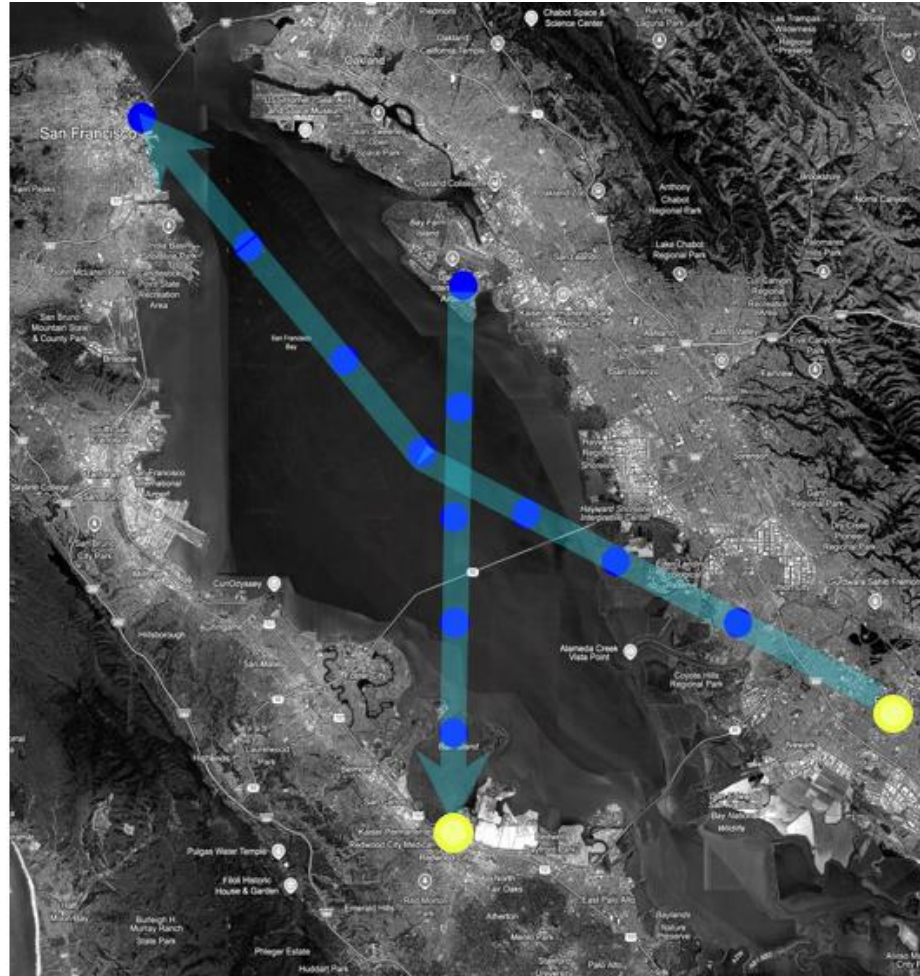
Safety-informed MARL



10x Speed


Air Taxi Operation Scenario: Intersecting corridors

Intersection Scenario




8 vehicles in each air corridor (total 16), crossing each other

Urban Air Mobility Simulation Result: Intersection Scenario



Avg. Travel Time: **781 sec**
Near Collision: 0.129%

Active CBVF-based
safety filtering



Avg. Travel Time: **660.8 sec**
Near Collision: **0.056%**

Resolving Conflicting Constraints in MARL with Layered Safety

- **Layered architecture:** Combines CBVF safety filtering with MARL policy for safe, efficient navigation.
- **Proven performance:** Faster travel, more waypoints covered, fewer conflicts.
- **Broad evaluation:** Tested on quadrotor and fixed-wing dynamics in complex scenarios.

Conclusions


Implications of this work:

- ✓ Our method integrates model-based safety tools from control theory (CBVFs) with learning-based methods (MARL),
- ✓ Together forming a framework that addresses two major challenges in multi-agent problems—safety and efficient coordination.
- ✓ Our work demonstrate the viability of hybrid approaches that combine learning and control and illustrate how RL can be responsibly applied in safety-critical settings.

Proposed Layered Safety Architecture

Multi-agent reinforcement learning




- Learns to proactively navigate potentially conflicting constraints
- Optimizes task performance (e.g., go to goal)



Control Barrier Value Functions (CBVFs)

a CBF design method based on Hamilton-Jacobi reachability

- Filters unsafe action generated by MARL
- Provides a safety-informed reward signal




Layered safety architecture

enhances both safety and performance

For more details check out our RSS 2025 Paper!

Resolving Conflicting Constraints in Multi-Agent Reinforcement Learning with Layered Safety

Jason J. Choi*, Jasmine Jerry Aloor*, Jingqi Li*, Maria G. Mendoza,
Hamsa Balakrishnan, Claire J. Tomlin



Webpage,
Paper, Videos

